

Predict Stock Prices using Time Series Analysis by Patrick BENIE

April 6, 2021

```
[1]: import pandas as pd
import numpy as np
```

```
[2]: df=pd.read_csv('/Users/patricksllearningprograms/Desktop/Python Projects/Stock_
↪Price Predictor/BAJFINANCE.csv')
df.head()
```

```
[2]:
```


	Date	Symbol	Series	Prev	Close	Open	High	Low	Last	\
0	2000-01-03	BAJAUTOFIN	EQ		46.95	49.45	50.75	46.5	50.75	
1	2000-01-04	BAJAUTOFIN	EQ		50.75	53.20	53.20	47.9	48.00	
2	2000-01-05	BAJAUTOFIN	EQ		48.10	46.55	47.40	44.6	44.60	
3	2000-01-06	BAJAUTOFIN	EQ		44.60	43.50	46.00	42.1	46.00	
4	2000-01-07	BAJAUTOFIN	EQ		45.25	48.00	48.00	42.0	42.90	

	Close	VWAP	Volume	Turnover	Trades	Deliverable	Volume	%Deliverble
0	50.75	50.05	7600	3.803800e+10	NaN		NaN	NaN
1	48.10	48.56	5000	2.428000e+10	NaN		NaN	NaN
2	44.60	45.47	3500	1.591450e+10	NaN		NaN	NaN
3	45.25	44.43	6200	2.754750e+10	NaN		NaN	NaN
4	42.90	44.44	3500	1.555550e+10	NaN		NaN	NaN

```
[3]: #since it's a time series use case, everything that happens in the data must be_
↪related to the time series index
#assign the 'Date' column to the time series index
df.set_index('Date',inplace=True) #inplace parameter used to update dataframe
```

```
[4]: #VWAP volume weighed average price
#check the trend of VWAP
#Plotting the target variable VWAP over time
df['VWAP'].plot()
```

```
[4]: <AxesSubplot:xlabel='Date'>
```



```
[ ]:
```

```
[5]: #Data processing
```

```
[6]: df.shape
```

```
[6]: (5070, 14)
```

```
[7]: #checking for missing values
df.isna().sum()
```

```
[7]: Symbol          0
      Series         0
      Prev Close     0
      Open          0
      High          0
      Low           0
      Last          0
      Close         0
      VWAP          0
      Volume        0
      Turnover       0
      Trades        2779
      Deliverable Volume  446
```

```
%Deliverble          446
dtype: int64
```

```
[8]: df.dropna(inplace=True)
```

```
[9]: df.isna().sum()
```

```
[9]: Symbol          0
     Series          0
     Prev Close      0
     Open            0
     High            0
     Low             0
     Last            0
     Close           0
     VWAP            0
     Volume          0
     Turnover        0
     Trades          0
     Deliverable Volume 0
     %Deliverble     0
     dtype: int64
```

```
[10]: df.shape
```

```
[10]: (2291, 14)
```

```
[11]: #create a copy of df
     data=df.copy()
```

```
[12]: data.dtypes
```

```
[12]: Symbol          object
     Series          object
     Prev Close      float64
     Open            float64
     High            float64
     Low             float64
     Last            float64
     Close           float64
     VWAP            float64
     Volume          int64
     Turnover        float64
     Trades          float64
     Deliverable Volume float64
     %Deliverble     float64
     dtype: object
```

```
[13]: data.columns
```

```
[13]: Index(['Symbol', 'Series', 'Prev Close', 'Open', 'High', 'Low', 'Last',  
        'Close', 'VWAP', 'Volume', 'Turnover', 'Trades', 'Deliverable Volume',  
        '%Deliverble'],  
        dtype='object')
```

```
[14]: #we choose 5 features(columns) for training purposes  
lag_features=['High','Low','Volume','Turnover','Trades']  
window1=3  
window2=7
```

```
[15]: #mean  
for feature in lag_features:  
    data[feature+'rolling_mean_3']=data[feature].rolling(window=window1).mean()  
    data[feature+'rolling_mean_7']=data[feature].rolling(window=window2).mean()
```

```
[16]: #standardisation  
for feature in lag_features:  
    data[feature+'rolling_std_3']=data[feature].rolling(window=window1).std()  
    data[feature+'rolling_std_7']=data[feature].rolling(window=window2).std()
```

```
[17]: data.head()
```

```
[17]:
```

	Symbol	Series	Prev Close	Open	High	Low	Last	\
Date								
2011-06-01	BAJFINANCE	EQ	616.70	617.00	636.50	616.00	627.00	
2011-06-02	BAJFINANCE	EQ	631.85	625.00	638.90	620.00	634.00	
2011-06-03	BAJFINANCE	EQ	633.45	625.15	637.80	620.00	623.00	
2011-06-06	BAJFINANCE	EQ	625.00	620.00	641.00	611.35	611.35	
2011-06-07	BAJFINANCE	EQ	614.00	604.00	623.95	604.00	619.90	

	Close	VWAP	Volume	...	Highrolling_std_3	Highrolling_std_7	\
Date				...			
2011-06-01	631.85	627.01	6894	...	NaN	NaN	
2011-06-02	633.45	636.04	2769	...	NaN	NaN	
2011-06-03	625.00	625.09	51427	...	1.201388	NaN	
2011-06-06	614.00	616.03	5446	...	1.625833	NaN	
2011-06-07	619.15	617.73	5991	...	9.062422	NaN	

	Lowrolling_std_3	Lowrolling_std_7	Volumerolling_std_3	\
Date				
2011-06-01	NaN	NaN	NaN	
2011-06-02	NaN	NaN	NaN	
2011-06-03	2.309401	NaN	26980.871860	
2011-06-06	4.994080	NaN	27352.695339	
2011-06-07	8.008797	NaN	26391.221653	

	Volumerolling_std_7	Turnoverrolling_std_3	Turnoverrolling_std_7 \
Date			
2011-06-01	NaN	NaN	NaN
2011-06-02	NaN	NaN	NaN
2011-06-03	NaN	1.685217e+12	NaN
2011-06-06	NaN	1.710136e+12	NaN
2011-06-07	NaN	1.652377e+12	NaN

	Tradesrolling_std_3	Tradesrolling_std_7
Date		
2011-06-01	NaN	NaN
2011-06-02	NaN	NaN
2011-06-03	670.500559	NaN
2011-06-06	148.769396	NaN
2011-06-07	78.270897	NaN

[5 rows x 34 columns]

```
[18]: data.columns
```

```
[18]: Index(['Symbol', 'Series', 'Prev Close', 'Open', 'High', 'Low', 'Last',
          'Close', 'VWAP', 'Volume', 'Turnover', 'Trades', 'Deliverable Volume',
          '%Deliverble', 'Highrolling_mean_3', 'Highrolling_mean_7',
          'Lowrolling_mean_3', 'Lowrolling_mean_7', 'Volumerolling_mean_3',
          'Volumerolling_mean_7', 'Turnoverrolling_mean_3',
          'Turnoverrolling_mean_7', 'Tradesrolling_mean_3',
          'Tradesrolling_mean_7', 'Highrolling_std_3', 'Highrolling_std_7',
          'Lowrolling_std_3', 'Lowrolling_std_7', 'Volumerolling_std_3',
          'Volumerolling_std_7', 'Turnoverrolling_std_3', 'Turnoverrolling_std_7',
          'Tradesrolling_std_3', 'Tradesrolling_std_7'],
          dtype='object')
```

```
[19]: data.shape
```

```
[19]: (2291, 34)
```

```
[20]: data.isna().sum()
```

```
[20]: Symbol          0
      Series          0
      Prev Close      0
      Open            0
      High            0
      Low             0
      Last            0
      Close           0
```

VWAP	0
Volume	0
Turnover	0
Trades	0
Deliverable Volume	0
%Deliverble	0
Highrolling_mean_3	2
Highrolling_mean_7	6
Lowrolling_mean_3	2
Lowrolling_mean_7	6
Volumerolling_mean_3	2
Volumerolling_mean_7	6
Turnoverrolling_mean_3	2
Turnoverrolling_mean_7	6
Tradesrolling_mean_3	2
Tradesrolling_mean_7	6
Highrolling_std_3	2
Highrolling_std_7	6
Lowrolling_std_3	2
Lowrolling_std_7	6
Volumerolling_std_3	2
Volumerolling_std_7	6
Turnoverrolling_std_3	2
Turnoverrolling_std_7	6
Tradesrolling_std_3	2
Tradesrolling_std_7	6

dtype: int64

```
[21]: data.dropna(inplace=True)
```

```
[22]: data.isna().sum()
```

```
[22]: Symbol      0
      Series      0
      Prev Close  0
      Open        0
      High        0
      Low         0
      Last        0
      Close       0
      VWAP        0
      Volume      0
      Turnover    0
      Trades      0
      Deliverable Volume  0
      %Deliverble  0
      Highrolling_mean_3  0
```

```

Highrolling_mean_7      0
Lowrolling_mean_3       0
Lowrolling_mean_7       0
Volumerolling_mean_3    0
Volumerolling_mean_7    0
Turnoverrolling_mean_3  0
Turnoverrolling_mean_7  0
Tradesrolling_mean_3    0
Tradesrolling_mean_7    0
Highrolling_std_3       0
Highrolling_std_7       0
Lowrolling_std_3        0
Lowrolling_std_7        0
Volumerolling_std_3     0
Volumerolling_std_7     0
Turnoverrolling_std_3   0
Turnoverrolling_std_7   0
Tradesrolling_std_3     0
Tradesrolling_std_7     0
dtype: int64

```

```
[23]: data.columns
```

```

[23]: Index(['Symbol', 'Series', 'Prev Close', 'Open', 'High', 'Low', 'Last',
            'Close', 'VWAP', 'Volume', 'Turnover', 'Trades', 'Deliverable Volume',
            '%Deliverble', 'Highrolling_mean_3', 'Highrolling_mean_7',
            'Lowrolling_mean_3', 'Lowrolling_mean_7', 'Volumerolling_mean_3',
            'Volumerolling_mean_7', 'Turnoverrolling_mean_3',
            'Turnoverrolling_mean_7', 'Tradesrolling_mean_3',
            'Tradesrolling_mean_7', 'Highrolling_std_3', 'Highrolling_std_7',
            'Lowrolling_std_3', 'Lowrolling_std_7', 'Volumerolling_std_3',
            'Volumerolling_std_7', 'Turnoverrolling_std_3', 'Turnoverrolling_std_7',
            'Tradesrolling_std_3', 'Tradesrolling_std_7'],
            dtype='object')

```

```

[24]: #independent features
ind_features=['Highrolling_mean_3', 'Highrolling_mean_7',
             'Lowrolling_mean_3', 'Lowrolling_mean_7', 'Volumerolling_mean_3',
             'Volumerolling_mean_7', 'Turnoverrolling_mean_3',
             'Turnoverrolling_mean_7', 'Tradesrolling_mean_3',
             'Tradesrolling_mean_7', 'Highrolling_std_3', 'Highrolling_std_7',
             'Lowrolling_std_3', 'Lowrolling_std_7', 'Volumerolling_std_3',
             'Volumerolling_std_7', 'Turnoverrolling_std_3', 'Turnoverrolling_std_7',
             'Tradesrolling_std_3', 'Tradesrolling_std_7']

```

```

[25]: training_data=data[0:1800] # about 80% of the dataset
      test_data=data[1800:]

```

[26]: training_data

[26]:

	Symbol	Series	Prev Close	Open	High	Low	Last \
Date							
2011-06-09	BAJFINANCE	EQ	635.60	639.80	647.00	630.00	630.00
2011-06-10	BAJFINANCE	EQ	631.10	641.85	648.25	618.55	621.10
2011-06-13	BAJFINANCE	EQ	622.20	616.00	627.85	616.00	622.75
2011-06-14	BAJFINANCE	EQ	624.95	625.00	628.95	619.95	621.20
2011-06-15	BAJFINANCE	EQ	622.10	612.00	623.00	598.10	605.00
...
2018-09-04	BAJFINANCE	EQ	2724.05	2724.00	2777.65	2683.50	2748.00
2018-09-05	BAJFINANCE	EQ	2746.30	2740.15	2764.80	2668.00	2704.45
2018-09-06	BAJFINANCE	EQ	2716.90	2729.00	2731.50	2671.40	2672.20
2018-09-07	BAJFINANCE	EQ	2684.10	2698.40	2751.40	2672.60	2745.00
2018-09-10	BAJFINANCE	EQ	2744.20	2732.00	2738.00	2596.00	2607.60

	Close	VWAP	Volume	...	Highrolling_std_3 \
Date				...	
2011-06-09	631.10	638.27	31252	...	12.769789
2011-06-10	622.20	634.16	30885	...	1.639360
2011-06-13	624.95	622.92	3981	...	11.434196
2011-06-14	622.10	625.35	5597	...	11.473593
2011-06-15	601.70	606.90	12590	...	3.165833
...
2018-09-04	2746.30	2726.23	2606992	...	88.954937
2018-09-05	2716.90	2712.53	1728455	...	63.129081
2018-09-06	2684.10	2695.89	1147879	...	23.818183
2018-09-07	2744.20	2716.32	1264436	...	16.755397
2018-09-10	2615.65	2655.39	1570179	...	10.147413

	Highrolling_std_7	Lowrolling_std_3	Lowrolling_std_7 \
Date			
2011-06-09	7.494911	15.011107	9.410145
2011-06-10	8.227994	13.030765	9.501961
2011-06-13	9.497080	7.456597	9.298317
2011-06-14	10.198891	2.002707	9.293713
2011-06-15	11.352292	11.643560	11.262712
...
2018-09-04	79.489416	83.341306	107.041856
2018-09-05	100.594924	24.113551	119.854378
2018-09-06	113.135709	8.146778	118.187686
2018-09-07	106.101111	2.386071	100.988340
2018-09-10	84.670766	43.882722	74.650851

	Volumerolling_std_3	Volumerolling_std_7	Turnoverrolling_std_3 \
Date			
2011-06-09	13497.047986	18373.894011	8.665128e+11

2011-06-10	2434.970705	18047.331029	1.579562e+11
2011-06-13	15640.051929	17835.642665	9.982094e+11
2011-06-14	15088.183102	13139.472184	9.595224e+11
2011-06-15	4575.812970	12469.137006	2.733211e+11
...
2018-09-04	696998.737787	642979.884199	1.827289e+14
2018-09-05	483305.554092	570427.041296	1.375623e+14
2018-09-06	734609.476628	584307.816124	2.020435e+14
2018-09-07	307128.544854	599913.723573	8.394860e+13
2018-09-10	218098.451125	592650.415369	5.493903e+13

Date	Turnoverrolling_std_7	Tradesrolling_std_3	Tradesrolling_std_7
2011-06-09	1.155773e+12	354.841279	463.942320
2011-06-10	1.136003e+12	272.875429	448.020620
2011-06-13	1.123423e+12	557.373603	464.778596
2011-06-14	8.397765e+11	656.385050	455.234163
2011-06-15	7.983186e+11	150.639747	436.941971
...
2018-09-04	1.673270e+14	2794.744413	21708.856675
2018-09-05	1.467744e+14	21770.352164	17267.163548
2018-09-06	1.561692e+14	30923.312263	24013.596897
2018-09-07	1.641199e+14	12123.301421	25604.694191
2018-09-10	1.647685e+14	15501.359822	25567.783116

[1800 rows x 34 columns]

[]:

[27]: !pip install pmdarima

```
Requirement already satisfied: pmdarima in
/Applications/anaconda3/lib/python3.8/site-packages (1.8.0)
Requirement already satisfied: statsmodels!=0.12.0,>=0.11 in
/Applications/anaconda3/lib/python3.8/site-packages (from pmdarima) (0.12.2)
Requirement already satisfied: scikit-learn>=0.22 in
/Applications/anaconda3/lib/python3.8/site-packages (from pmdarima) (0.23.2)
Requirement already satisfied: setuptools!=50.0.0,>=38.6.0 in
/Applications/anaconda3/lib/python3.8/site-packages (from pmdarima)
(50.3.1.post20201107)
Requirement already satisfied: pandas>=0.19 in
/Applications/anaconda3/lib/python3.8/site-packages (from pmdarima) (1.1.3)
Requirement already satisfied: numpy>=1.17.3 in
/Applications/anaconda3/lib/python3.8/site-packages (from pmdarima) (1.19.2)
Requirement already satisfied: joblib>=0.11 in
/Applications/anaconda3/lib/python3.8/site-packages (from pmdarima) (0.17.0)
Requirement already satisfied: scipy>=1.3.2 in
```

```

/Applications/anaconda3/lib/python3.8/site-packages (from pmdarima) (1.5.2)
Requirement already satisfied: urllib3 in
/Applications/anaconda3/lib/python3.8/site-packages (from pmdarima) (1.25.11)
Requirement already satisfied: Cython<0.29.18,>=0.29 in
/Applications/anaconda3/lib/python3.8/site-packages (from pmdarima) (0.29.17)
Requirement already satisfied: patsy>=0.5 in
/Applications/anaconda3/lib/python3.8/site-packages (from
statsmodels!=0.12.0,>=0.11->pmdarima) (0.5.1)
Requirement already satisfied: threadpoolctl>=2.0.0 in
/Applications/anaconda3/lib/python3.8/site-packages (from scikit-
learn>=0.22->pmdarima) (2.1.0)
Requirement already satisfied: pytz>=2017.2 in
/Applications/anaconda3/lib/python3.8/site-packages (from
pandas>=0.19->pmdarima) (2020.1)
Requirement already satisfied: python-dateutil>=2.7.3 in
/Applications/anaconda3/lib/python3.8/site-packages (from
pandas>=0.19->pmdarima) (2.8.1)
Requirement already satisfied: six in
/Applications/anaconda3/lib/python3.8/site-packages (from
patsy>=0.5->statsmodels!=0.12.0,>=0.11->pmdarima) (1.15.0)

```

```
[28]: from pmdarima import auto_arima
```

```
[29]: import warnings
warnings.filterwarnings('ignore')
```

```
[30]: #find the best arima model
model=auto_arima(y=training_data['VWAP'],exogenous=training_data[ind_features],trace=True)
```

```

Performing stepwise search to minimize aic
ARIMA(2,0,2)(0,0,0)[0] intercept : AIC=20931.536, Time=4.79 sec
ARIMA(0,0,0)(0,0,0)[0] intercept : AIC=20925.224, Time=2.57 sec
ARIMA(1,0,0)(0,0,0)[0] intercept : AIC=20926.348, Time=2.35 sec
ARIMA(0,0,1)(0,0,0)[0] intercept : AIC=20926.320, Time=3.12 sec
ARIMA(0,0,0)(0,0,0)[0]          : AIC=32616.913, Time=2.13 sec
ARIMA(1,0,1)(0,0,0)[0] intercept : AIC=20929.234, Time=3.25 sec

```

```

Best model: ARIMA(0,0,0)(0,0,0)[0] intercept
Total fit time: 18.224 seconds

```

```
[31]: #Arima model
model.fit(training_data['VWAP'],training_data[ind_features])
```


```
[31]: ARIMA(order=(0, 0, 0), scoring_args={}, suppress_warnings=True)
```

```
[32]: forecast=model.predict(n_periods=len(test_data),
    ↪exogenous=test_data[ind_features])
```

```
[33]: test_data['Forecast_ARIMA']=forecast
```

```
[34]: #blue line is the actual data and orange line is the prediction. we can see  
      ↳ that they're very close to each other  
      test_data[['VWAP','Forecast_ARIMA']].plot(figsize=(14,7))
```

```
[34]: <AxesSubplot:xlabel='Date'>
```



```
[ ]:
```

```
[35]: from sklearn.metrics import mean_absolute_error, mean_squared_error
```

```
[36]: np.sqrt(mean_squared_error(test_data['VWAP'],test_data['Forecast_ARIMA']))
```

```
[36]: 187.77545472730867
```

```
[37]: mean_absolute_error(test_data['VWAP'],test_data['Forecast_ARIMA'])
```

```
[37]: 124.6480740444013
```

```
[ ]:
```